

# ML Orchestration

Robert Clements

MSDS Program

University of San Francisco




# What to Expect

- Goal: to learn about creating reproducible and scalable ML workflows.
- How: we will study ML workflows and practice ML Orchestration with Metaflow and GCP and Kubernetes.


# ML Workflows

# ML Workflows




# Why do we need ML workflow orchestration?

- ML systems are very complex with many interconnected complex parts:
  - Variety of data sources
  - Variety of models
  - Variety of decisions
  - Volume, variety, velocity
- ML systems evolve
- Resource-heavy



# ML Workflows



Versioning: save versions of every run of the flow

Metadata: save the metadata of every run of the flow

Artifacts: save artifacts of every run of the flow

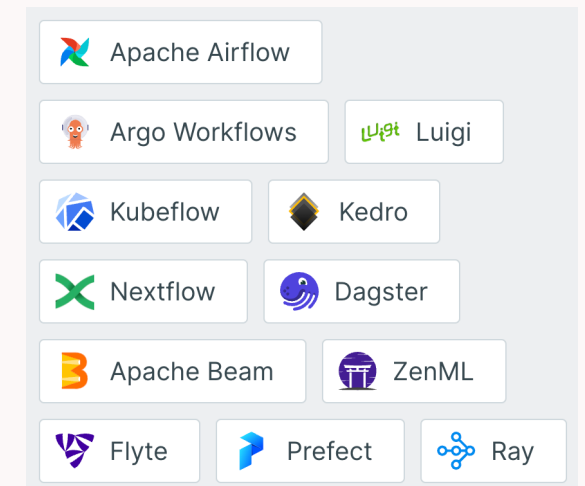
Compute: scale vertically and horizontally

Orchestration: run steps in order, parallel

Deployment: schedule/trigger the flow to run

# ML Workflows with Metaflow

- [Metaflow](#) began at Netflix in 2017 by Ville Tuulos
- Book - Effective Data Science Infrastructure
- Open-sourced soon thereafter
- DAG-based tool for creating scalable workflows, but focused on ML
- Runs on
  - Local machine – low-maintenance
  - AWS Batch – low-maintenance prototyping stack
  - AWS Batch + Step Functions – low-maintenance full stack
  - AWS/GCP/Azure K8s – customizable full stack
- Popular alternatives:
  - [Airflow](#)
  - [Flyte](#)
  - [Kubeflow](#)





# Why Metaflow?

- Designed for ML and for Data Scientists to use
- *Easy* to scale vertically and horizontally
- Reproducible and shareable workflows
- Covers the full stack:
  - Data
  - Compute
  - Orchestration
  - Versioning
  - Deployment
  - Modeling

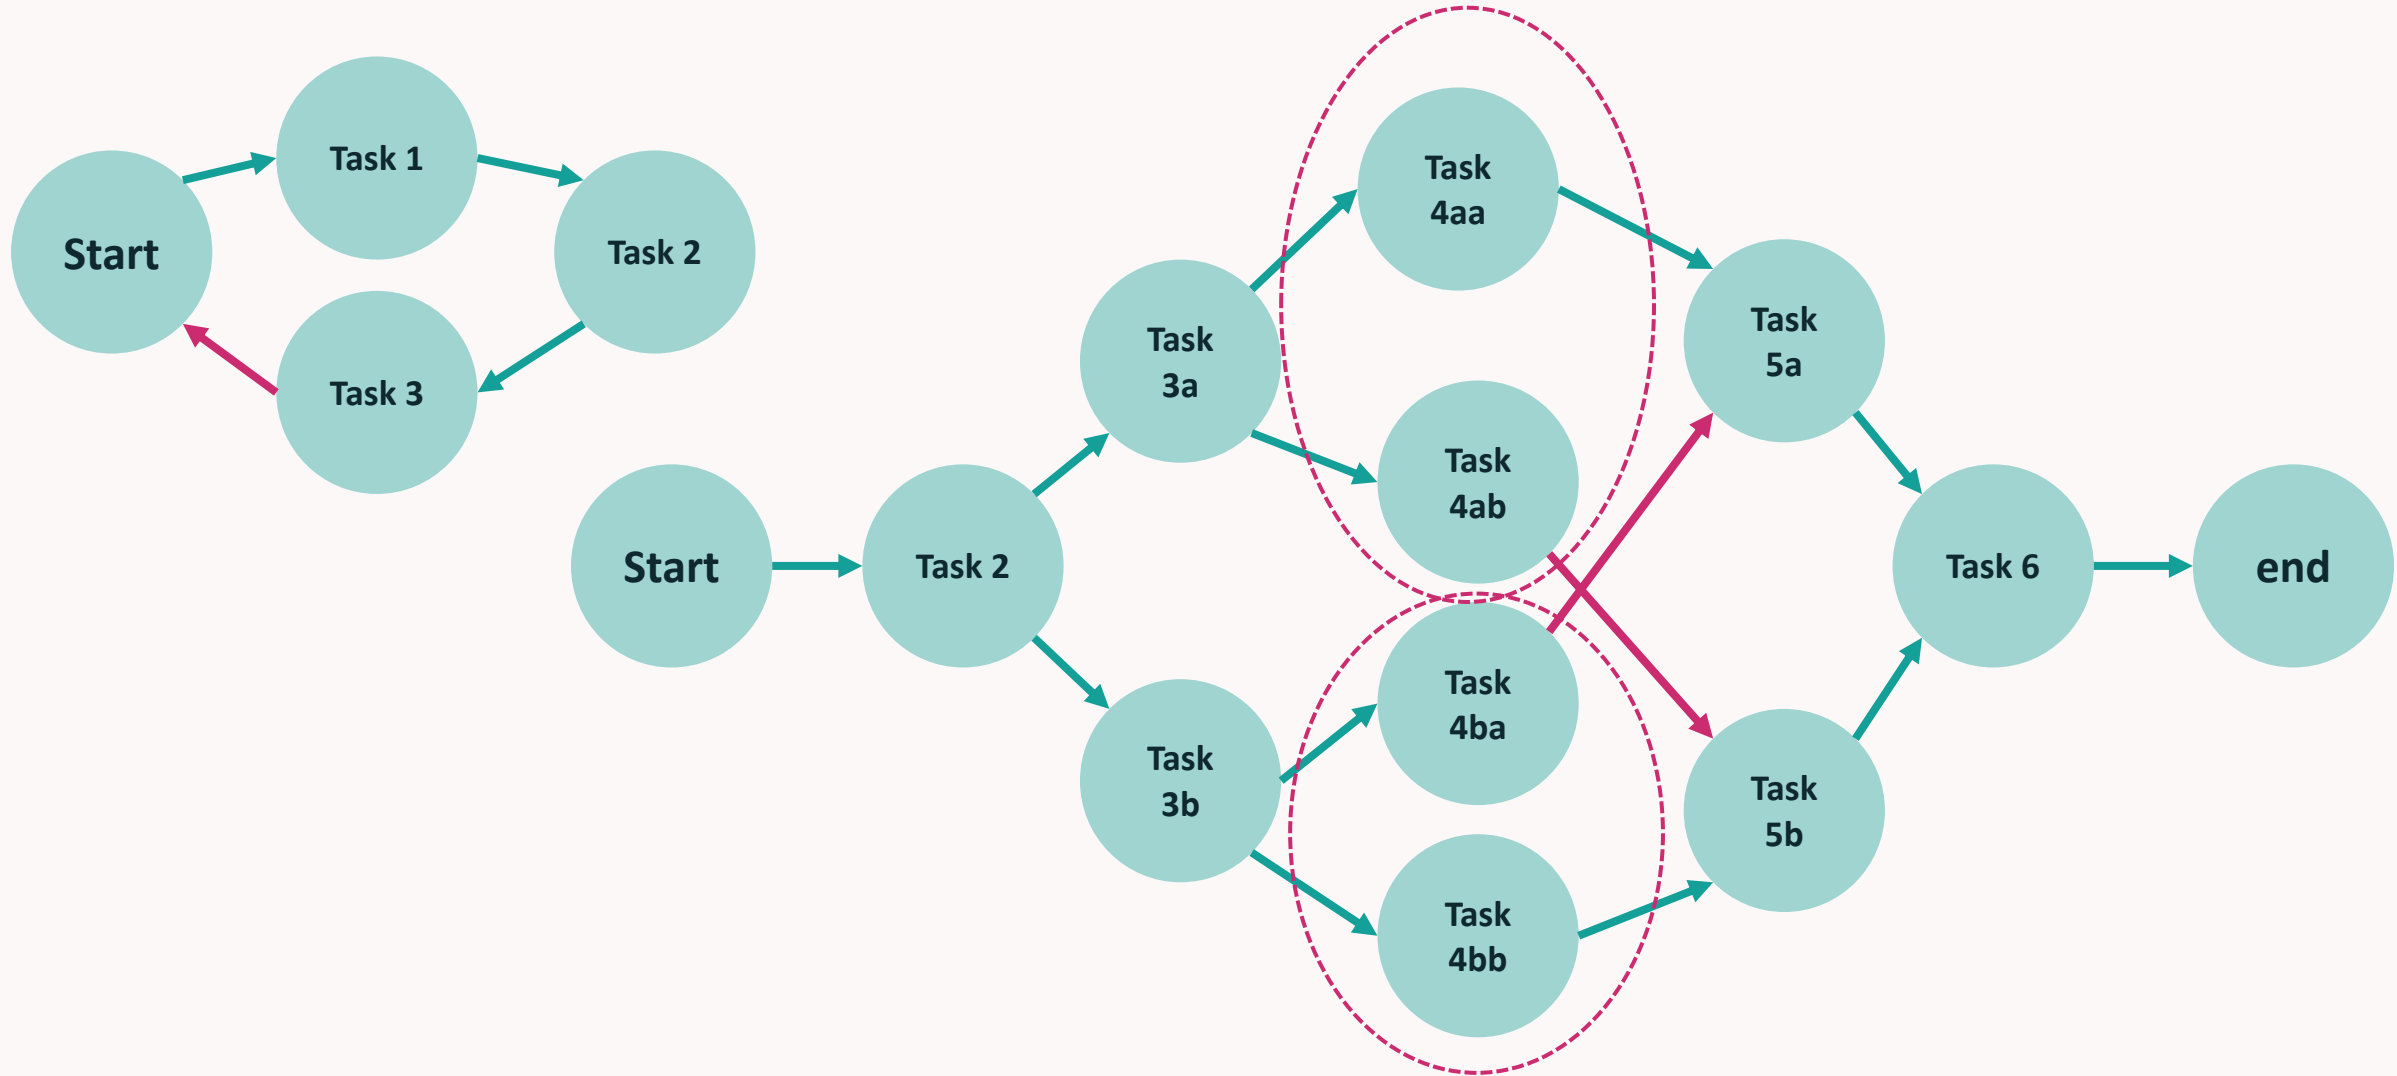




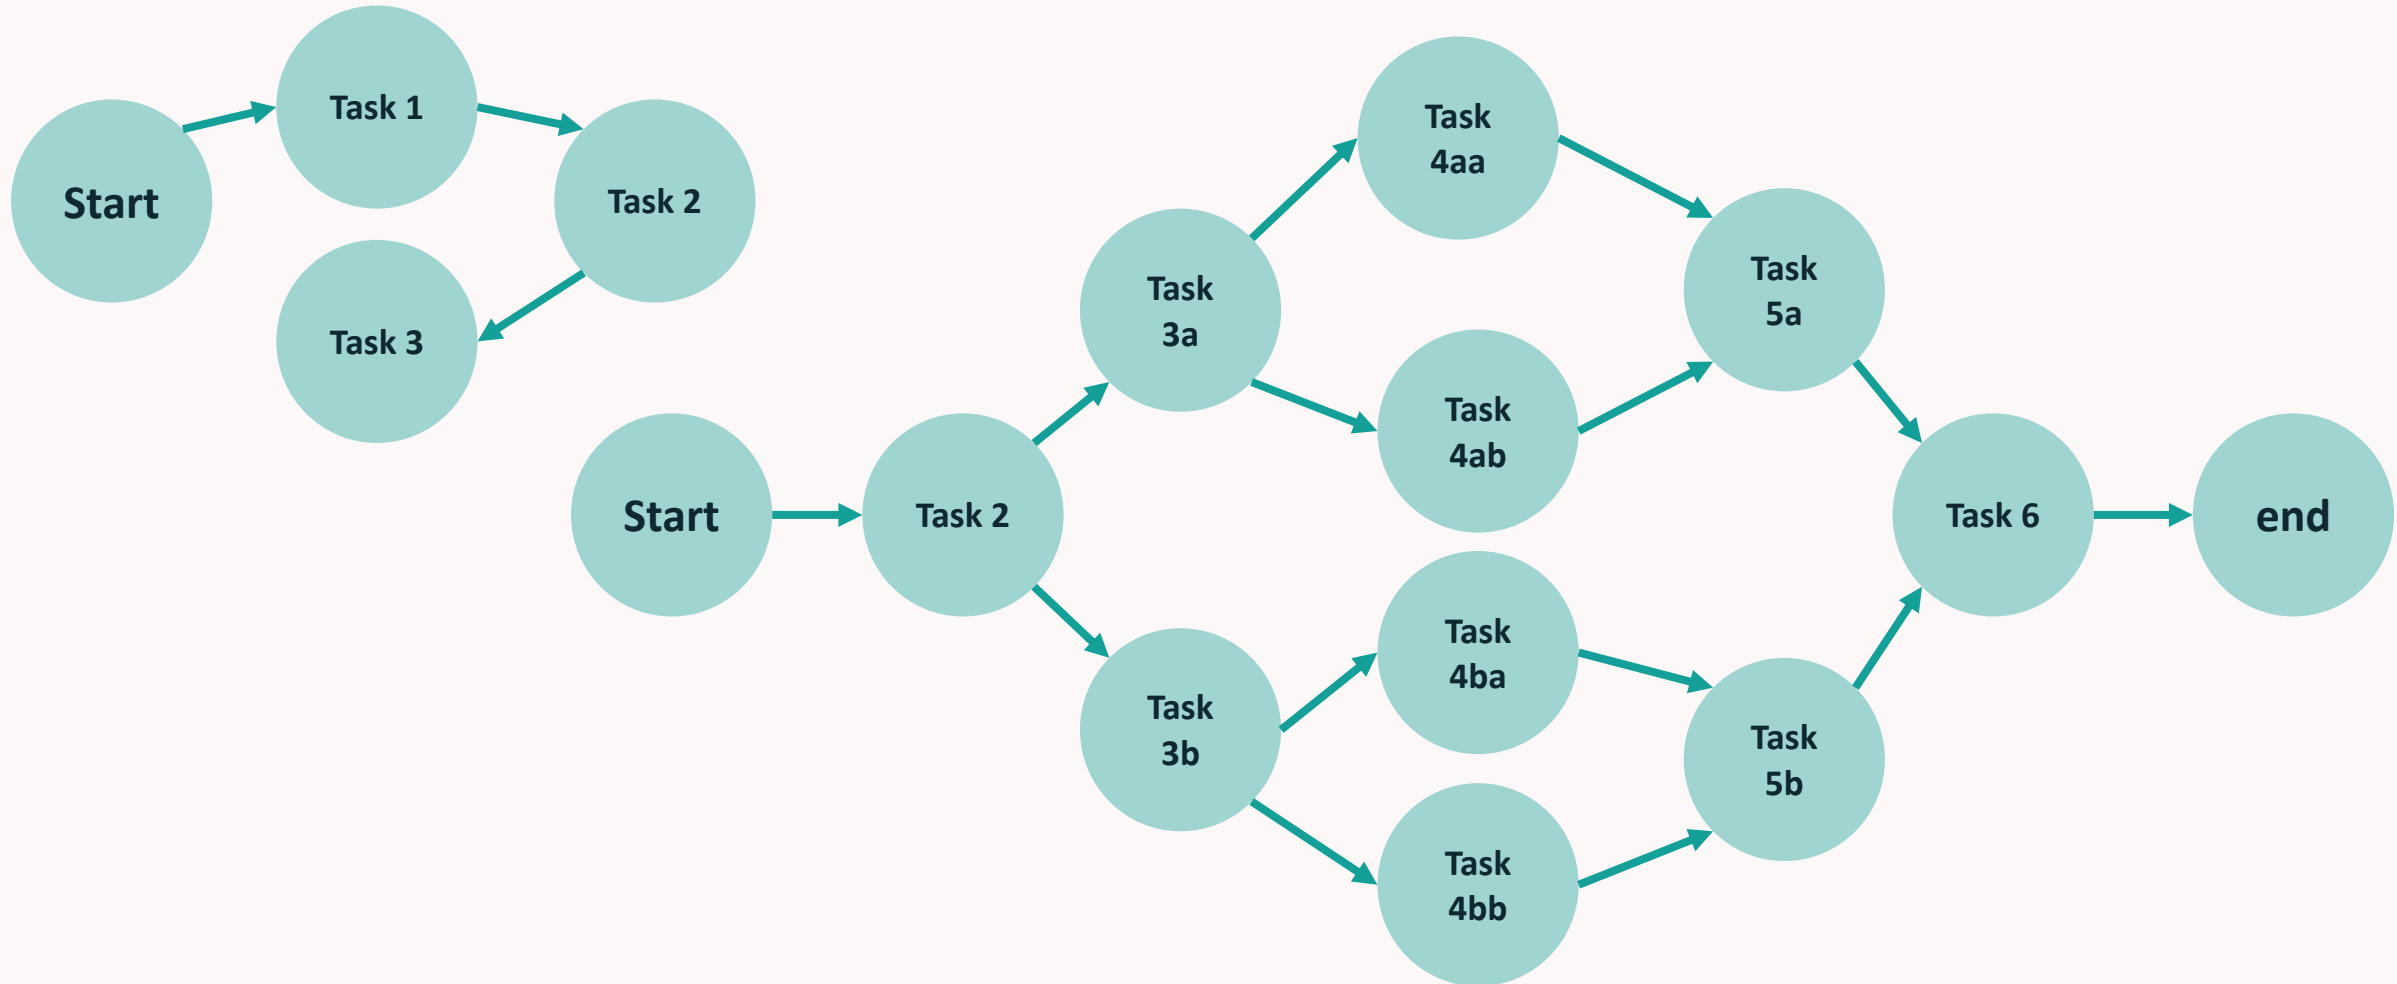
# ML Workflows – Valid DAGs



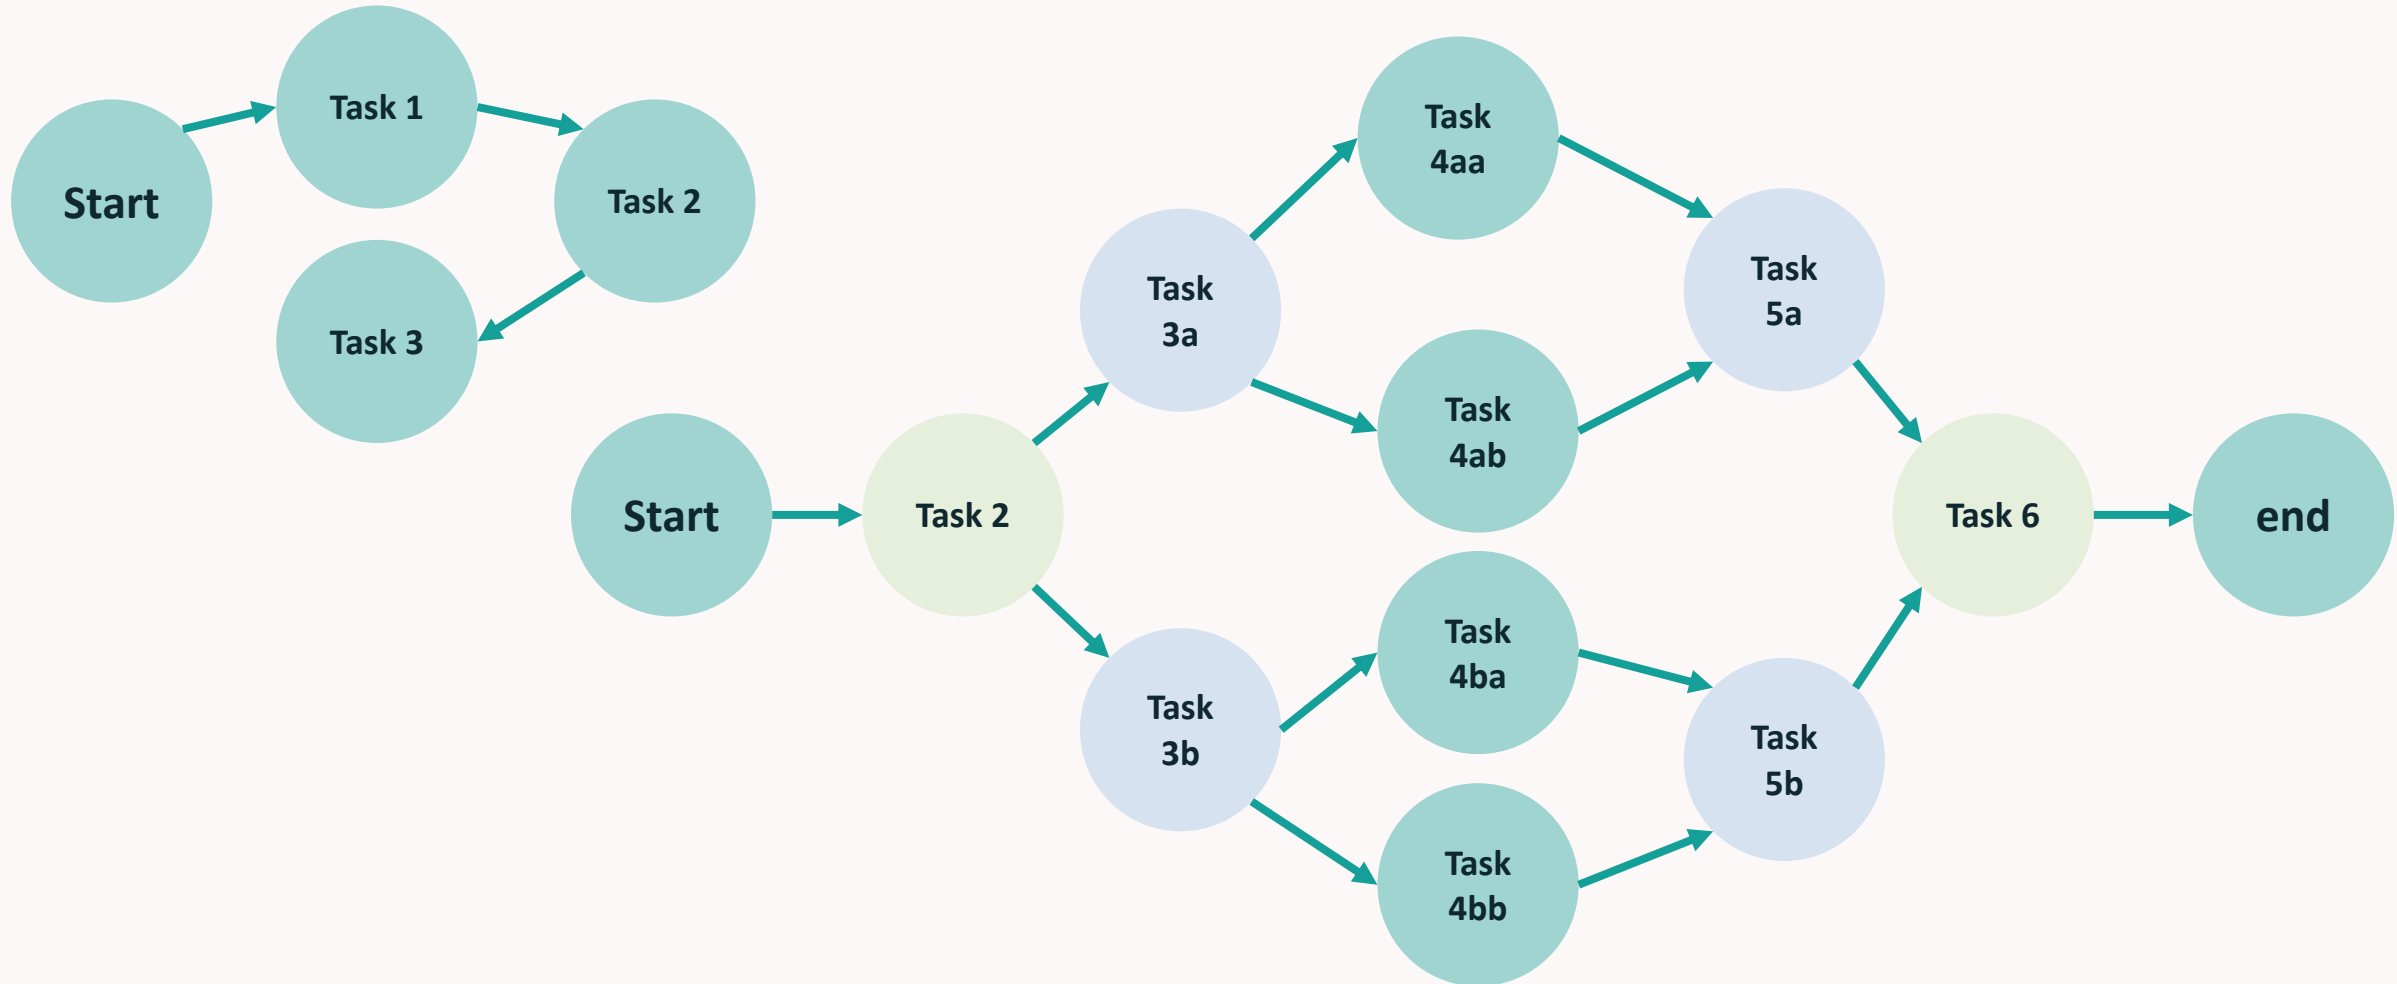
# ML Workflows - Invalid DAGs




# ML Workflows – Fixing Invalid DAGs



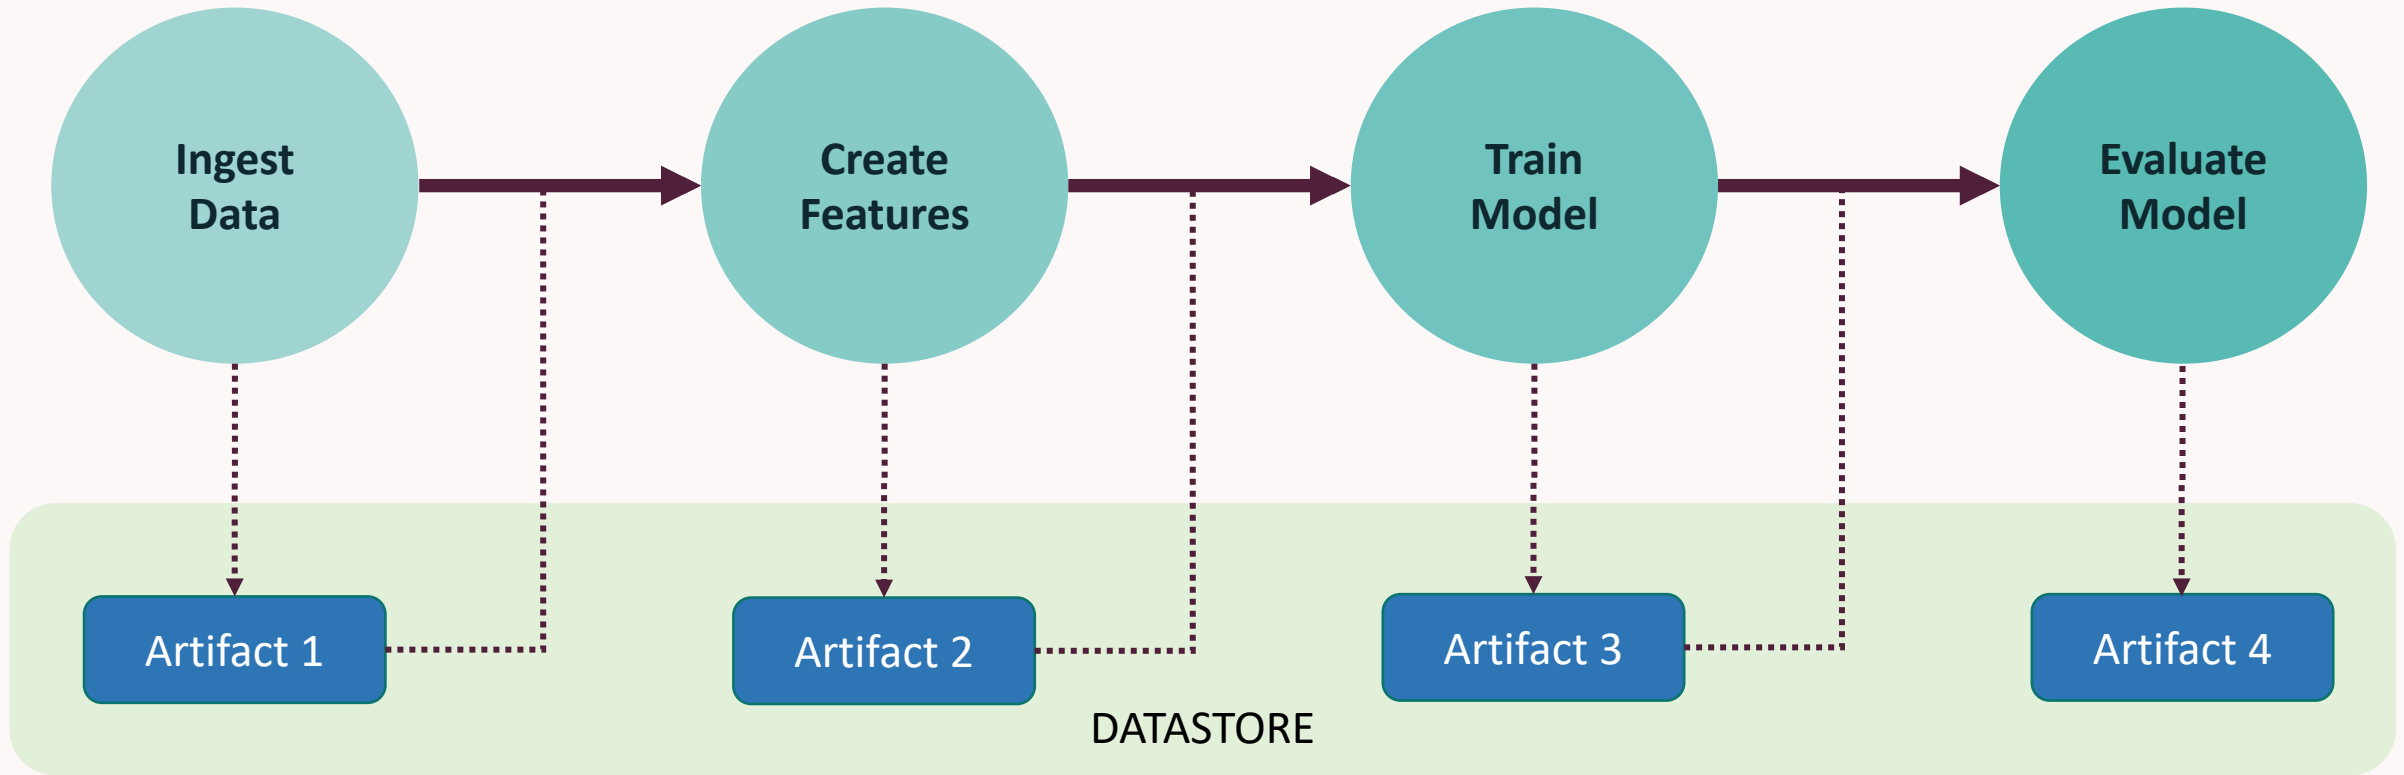
# ML Workflows – Fixing Invalid DAGs



# Static and Dynamic DAGs




# State Persistence




# Two Important Flows

## Model Training



## Model Scoring



# ML Pipeline Tools

- Metaflow
- Prefect
- Airflow
- Dagster
- ZenML
- Mage
- Kubeflow
- So many others...

How do you define the flow?

Where and how do you run the flow?



# ML Workflows in Metaflow

```
from metaflow import FlowSpec, step

class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)
```




Start

# ML Workflows in Metaflow

```
from metaflow import FlowSpec, step

class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)

    @step
    def one(self):
        ...
        self.next(self.two, self.three) # branching
```



# ML Workflows in Metaflow

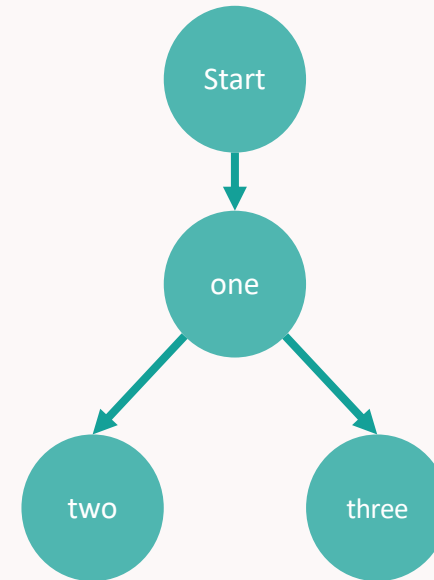
```
from metaflow import FlowSpec, step

class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)

    @step
    def one(self):
        ...
        self.next(self.two, self.three) # branching

    @step
    def two(self):
        ...
        self.next(self.four)

    @step
    def three(self):
        ...
        self.next(self.four)
```



# ML Workflows in Metaflow

```
from metaflow import FlowSpec, step


class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)

    @step
    def one(self):
        ...
        self.next(self.two, self.three) # branching

    @step
    def two(self):
        ...
        self.next(self.four)

    @step
    def three(self):
        ...
        self.next(self.four)

    @step
    def four(self, inputs): # join step
        ...
        self.next(self.end)
```



# ML Workflows in Metaflow

```
from metaflow import FlowSpec, step

class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)


    @step
    def one(self):
        ...
        self.next(self.two, self.three) # branching

    @step
    def two(self):
        ...
        self.next(self.four)

    @step
    def three(self):
        ...
        self.next(self.four)

    @step
    def four(self, inputs): # join step
        ...
        self.next(self.end)

    @step
    def end(self): # end of flow
        ...
```



# ML Workflows in Metaflow (dynamic)

```
from metaflow import FlowSpec, step


class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
        self.next(self.one)

    @step
    def one(self):
        iters = [list of things to iterate over]
        ...
        self.next(self.two, foreach='iters') # branching


    @step
    def two(self):
        # do something with self.input
        ...
        self.next(self.three)

    @step
    def three(self, inputs): # join step
        ...
        self.next(self.end)

    @step
    def end(self): # end of flow
        ...
```



# Scaling Vertically and Horizontally



# Scaling Vertically Using @resources

```
@resources(memory=60000, cpu=1)
@step
def start(self):
    import numpy
    import time
    big_matrix =
numpy.random.randn((80000, 80000))
    t = time.time()
    self.sum = numpy.sum(big_matrix)
    self.took = time.time() - t
    self.next(self.end)
```

- Set memory, cpu, gpu, or shared\_memory
- Does ***not*** inherently scale up, must be paired with a scalable compute layer, e.g. AWS Batch or K8s



# Scaling Horizontally Using @kubernetes

```
$ python BigSum.py run --with kubernetes
```

```
@kubernetes(memory=60000, cpu=1)
@step
def start(self):
    import numpy
    import time
    big_matrix =
numpy.random.randn((80000, 80000))
    t = time.time()
    self.sum = numpy.sum(big_matrix)
    self.took = time.time() - t
    self.next(self.end)
```

- Running on command line forces **all** steps to run in K8s

OR

- Replace @resources with @kubernetes (or @batch if using AWS) to run specific steps on specific compute layers
- Note: @kubernetes has additional arguments available

# Being Careful When Using @kubernetes

```
$ python BigSum.py run --with kubernetes --max-num-splits 100
```

- To avoid using too many resources with parallel jobs, use either
  - `--max-num-splits N`
  - `--max-workers N`

# Retry Steps Using @retry

```
$ python BigSum.py run --with retry
```

```
@retry
@step
def start(self):
    import numpy
    import time
    big_matrix =
numpy.random.randn((80000, 80000))
    t = time.time()
    self.sum = numpy.sum(big_matrix)
    self.took = time.time() - t
    self.next(self.end)
```

- For transient platform issues, use @retry
- Recommended when using remote compute layer
- To avoid retries for specific steps, use @retry(times=0)
- Can set number of retries and minutes\_between\_retries

# Catch Exceptions Using @catch

```
@step
def start(self):
    self.params = range(3)
    self.next(self.compute, foreach='params')

@catch(var='compute_failed')
@step
def compute(self):
    self.div = self.input
    self.x = 5 / self.div
    self.next(self.join)

@step
def join(self, inputs):
    for input in inputs:
        if input.compute_failed:
            print('compute failed for
parameter: %d' % input.div)
    self.next(self.end)
```

- Use @catch to catch exceptions that are not transient compute layer issues
- Code has to be rewritten to know how to handle exceptions
- var is optional

# Timeout Using @timeout

```
@timeout(seconds=5)
@step
def start(self):
    import numpy
    import time
    big_matrix =
numpy.random.randn((80000, 80000))
    t = time.time()
    self.sum = numpy.sum(big_matrix)
    self.took = time.time() - t
    self.next(self.end)
```

- Use the @timeout decorator to avoid stuck code

# Accessing Data in Metaflow

Metaflow has the `metaflow.S3` module for accessing S3 data, but when it comes to data, just follow best practices:

- Tip: keep data ***loading*** and data ***transformations*** separate
- Whenever possible, use Metaflow artifacts
  - Anything assigned to self will be persisted to subsequent steps
- Try to avoid importing *local* files
- Use larger instances for larger datasets
- Use parquet + Apache Arrow or numpy (not pandas) when possible

# Manage Dependencies with @conda

```
python LinearFlow.py --environment=conda run
```

## Step-level

```
@conda(libraries={"pandas": "0.22.0"})  
def fit_model(self):  
    ...
```

- Create conda environments for each step

## Flow-level

```
@conda_base(libraries={'numpy': '1.15.4'}, python='3.6.5')  
class LinearFlow(FlowSpec):  
    ...
```

- Create conda environment for entire flow
- Can be combined with step-level environments

# Metaflow ML Flows Demo



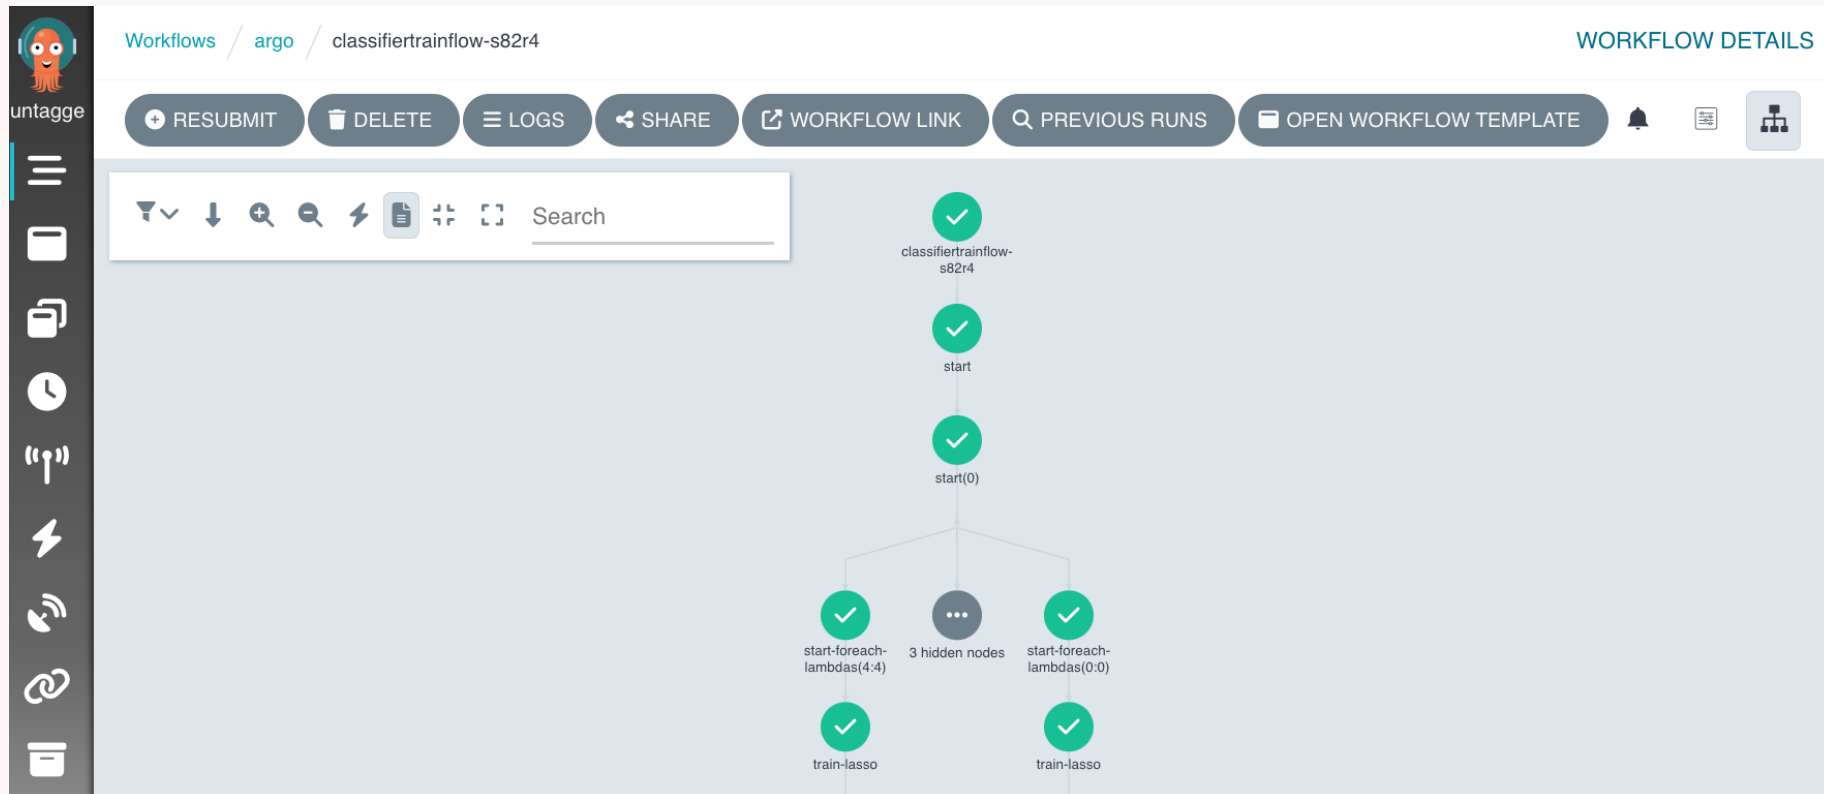
# Metaflow ML Flows Lab

# Metaflow in Production

- Use `@resources` and `@kubernetes` to scale
- Use `@conda` or `@conda_base` for packaging up the environment
- Use `@timeout`, `@retry`, and `@catch` to deal with problems
- You will need to:
  - Create a project in Google cloud
  - Install Gcloud CLI
  - Install Terraform
  - Install kubectl
  - Install Kubernetes, google-cloud-storage, google-auth python libraries
  - Patience and luck

# Scheduling Jobs with Argo Workflows

```
python LinearFlow.py --environment=conda --with retry argo-workflows create
```



# Scheduling Jobs with Argo Workflows

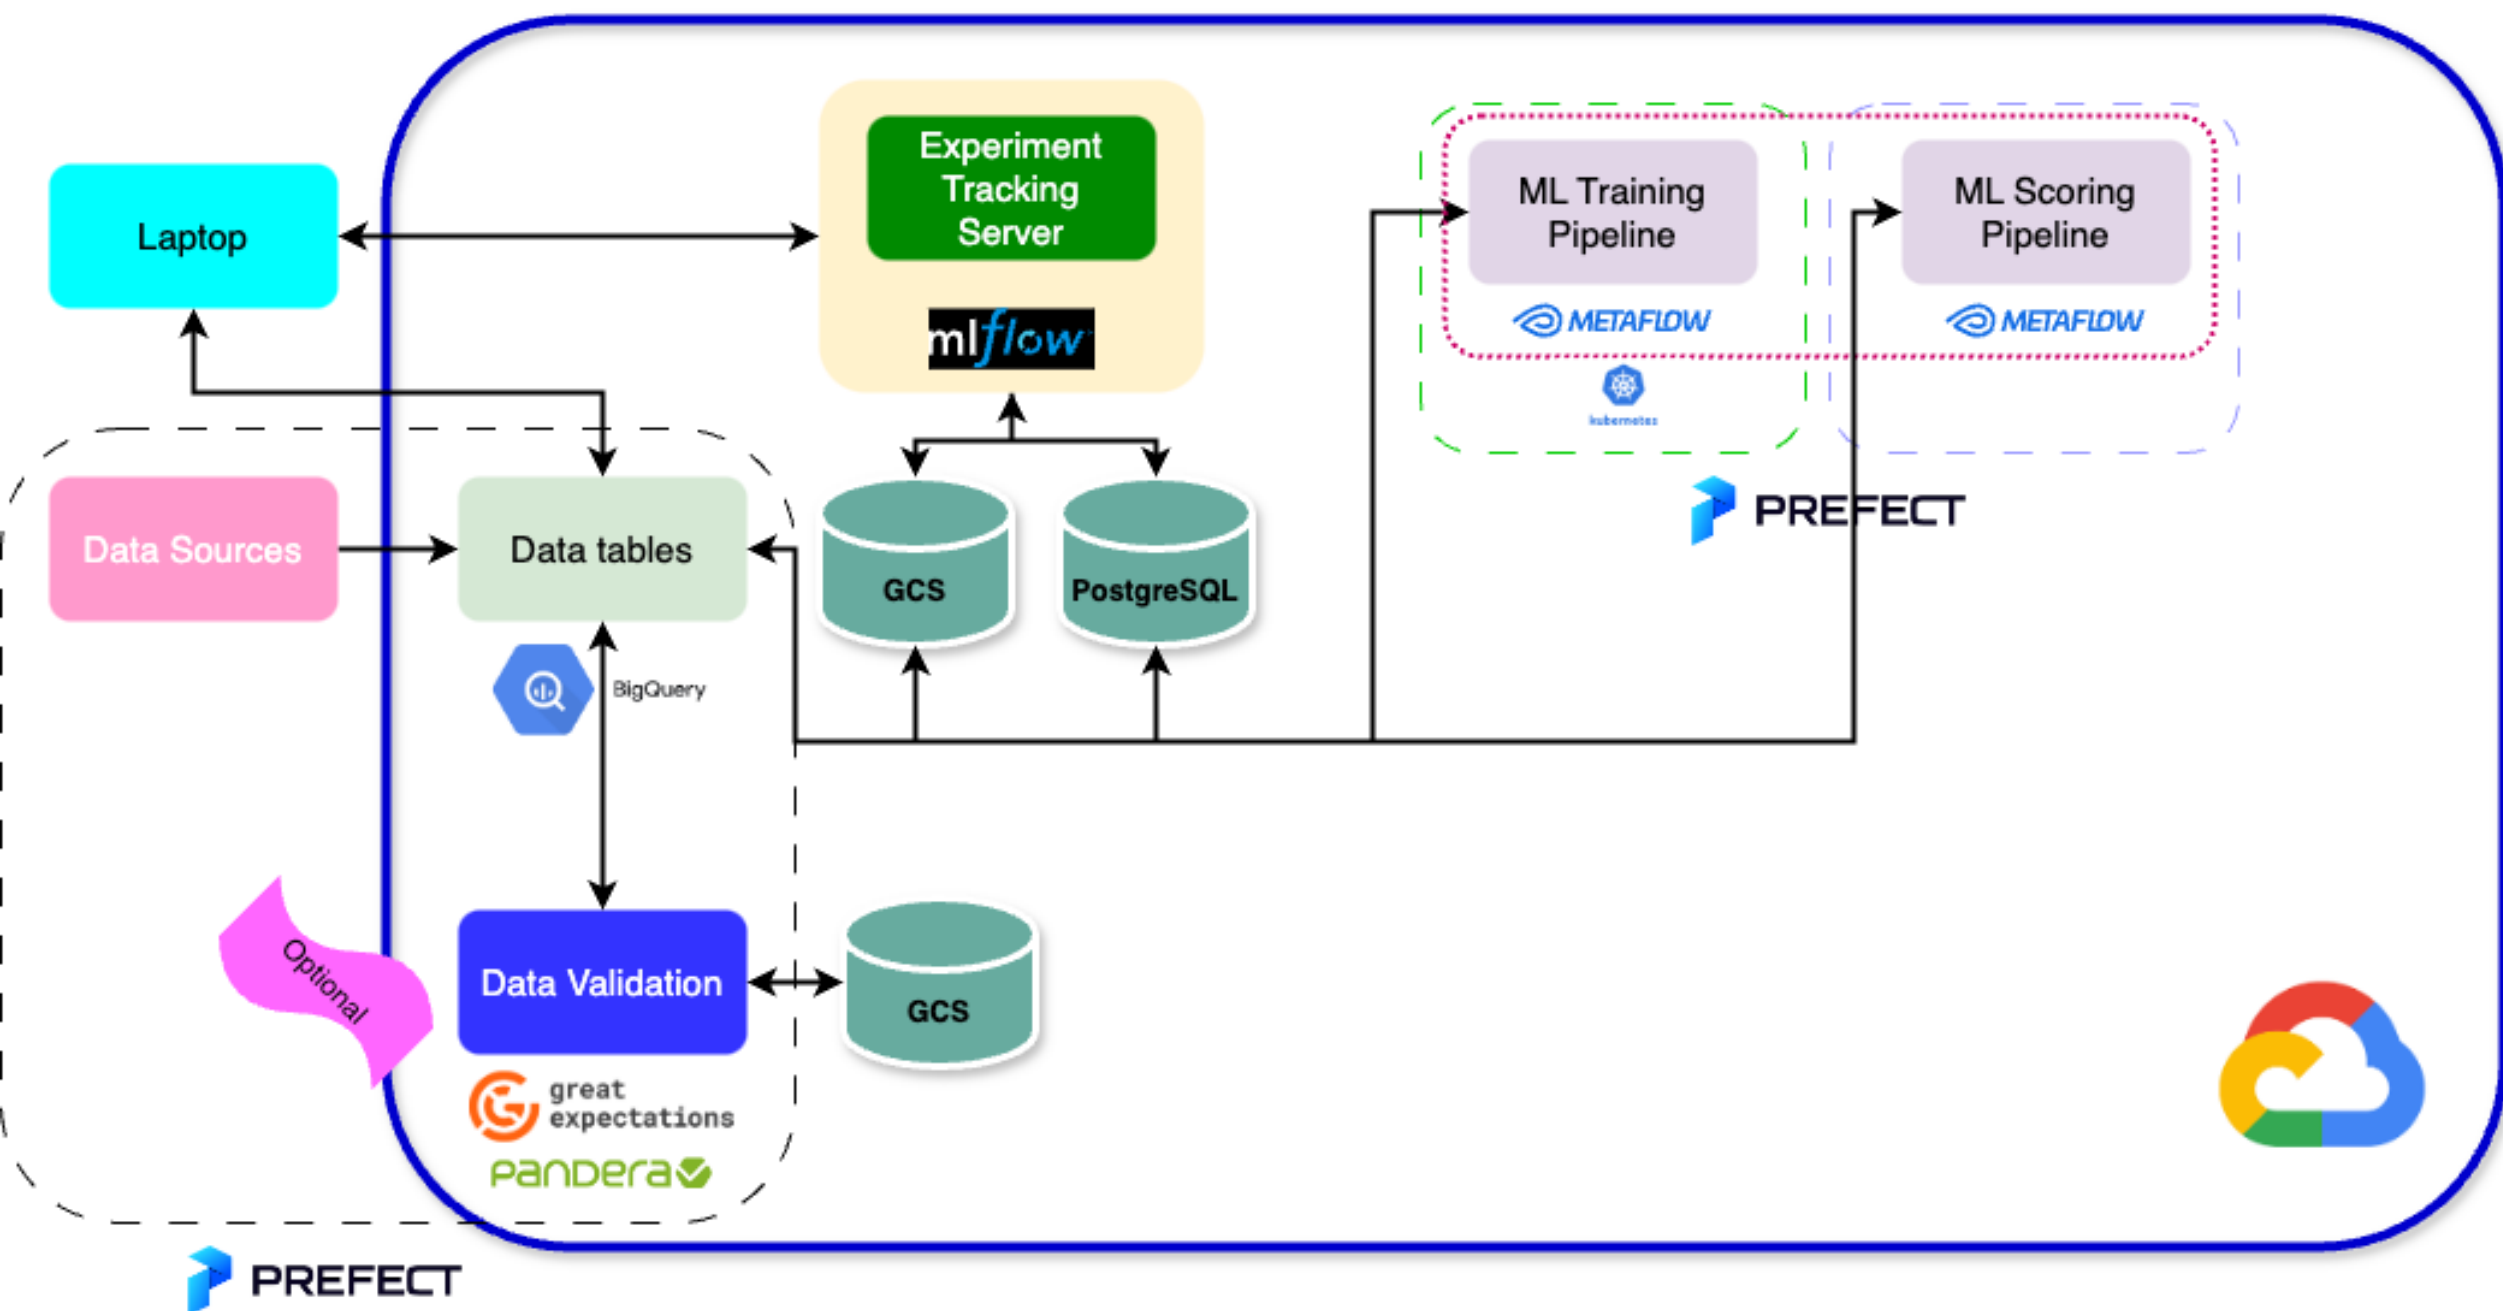
- Use @schedule decorator to schedule hourly, daily, weekly, or using a cron schedule
- Rename the flow (e.g. HelloWorldFlowStage) to deploy and test staging version before deploying production version

```
from metaflow import FlowSpec, step, schedule

@schedule(hourly=True)
class HelloWorldFlow(FlowSpec):
    @step
    def start(self): # start of flow
        ...
```

# More Things to Cover on Your Own – Read the Docs

- Inspecting Flows with the Client API
- Debugging Flows and the resume command
- Visualizing Results with Cards
- Organizing Flows with namespaces and tags



# Metaflow Scaling and Production Demo